ABHINAV TARIGOPPULA

Visakhapatnam, India | +91-9121611029 | abhinaaavvv07187@gmail.com

Linkedin | Github | Portfolio| Leetcode

PROFESSIONAL SUMMARY

Machine Learning Engineer with proven experience building and deploying production ML systems - from data pipeline engineering to real-time model serving and inference optimization. Experienced in LLM pipelines, RAG architectures, and end-to-end model lifecycle management using Python, LangChain, and FastAPI. Driven by measurable impact: improved model accuracy by 12% and reduced inference latency by 18% in production environments.

TECHNICAL SKILLS

Languages: Python, Java, JavaScript

ML & MODELLING: TensorFlow, PyTorch, Scikit-learn, LangChain, OpenAI API, NLP, Transformers, RAG, Generative AI

Backend & APIs: FastAPI, Node.js, REST APIs, Vector Databases (FAISS, Pinecone)

Frontend: React.js, Next.js

Cloud & DevOps: AWS, Docker, GitHub Actions, CI/CD

Tools: Git, SQL, Linux

WORK EXPERIENCE

PwC Launchpad Advisory Program (PwC India) | AI & Data Advisory Training ***Feb 2026 - Present***

Selected for PwC's flagship program focused on GenAI, Prompt Engineering, and Data Systems.

- Built and optimized LLM prompt workflows for real-world use cases.
- Gained hands-on exposure to enterprise data architectures and AI-driven solutions.
- Achieved Level 10 | 1500+ XP demonstrating top performance.

Machine Learning & AI Intern | OneStop AI ***May 2025 – Aug 2025 | Remote***

- Improved classification accuracy by 12% (from 78% to 90%) on 3 production models via hyperparameter tuning and architecture search.
- Cut real-time inference latency by 18% (from 220ms to 180ms) by refactoring preprocessing pipeline and applying model quantization.
- Engineered scalable end-to-end data pipelines covering preprocessing, feature engineering, and batch processing - integrated with CI/CD workflows for continuous model delivery.
- Strengthened model deployment reliability by integrating monitoring and performance tracking pipelines in collaboration with senior engineers, ensuring inference stability.

PROJECTS

PathForge | AI Placement Preparation System / *Pathforge*

Python, FastAPI, LangChain, OpenAI API, FAISS/Pinecone, React

- Architected an end-to-end AI placement preparation platform using LLM pipelines to deliver role-specific study material and mock interviews.
- Implemented Retrieval-Augmented Generation (RAG) with FAISS/Pinecone vector databases for context-aware Q&A and study material generation.
- Engineered intelligent roadmap generation engine analyzing user skill gaps to auto-generate structured learning paths
- Built LLM-driven modules for dynamic question generation, resume scoring, and mock interview simulation.
- Applied systematic prompt engineering to improve response accuracy by approximately 30% and reduce average token consumption per query.

MarkMe | MarkMe | Smart Attendance/ *kratos0718/MarkMe*

Python/JS, Face Recognition, GPS Geo-fencing, Liveness Detection

- Engineered a triple-layer attendance verification system combining real-time face recognition, GPS geo-fencing (100m radius), and rotating 6-digit session keys - eliminating proxy attendance.
- Implemented client-side face detection to ensure biometric data never leaves the user's device, addressing privacy constraints.
- Deployed live at GITAM University with session-based key expiry and liveness detection to block photo spoofing attacks.

EDUCATION

GITAM University, Andhra Pradesh **AUG 2023 – 2027**

B.Tech in Computer Science & Engineering (AI/ML Specialization) | CGPA: 8.04

FIITJEE Junior College

Intermediate(MPC)